

1.-Libraries

DONE BY BENNYPRIYA D

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as tlp
%matplotlib inline
import seaborn as ss
```

2.Loading the dataset

```
from google.colab import files
upload=files.upload()

a=pd.read_csv('/content/abalone.csv')

a.head()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight \
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395

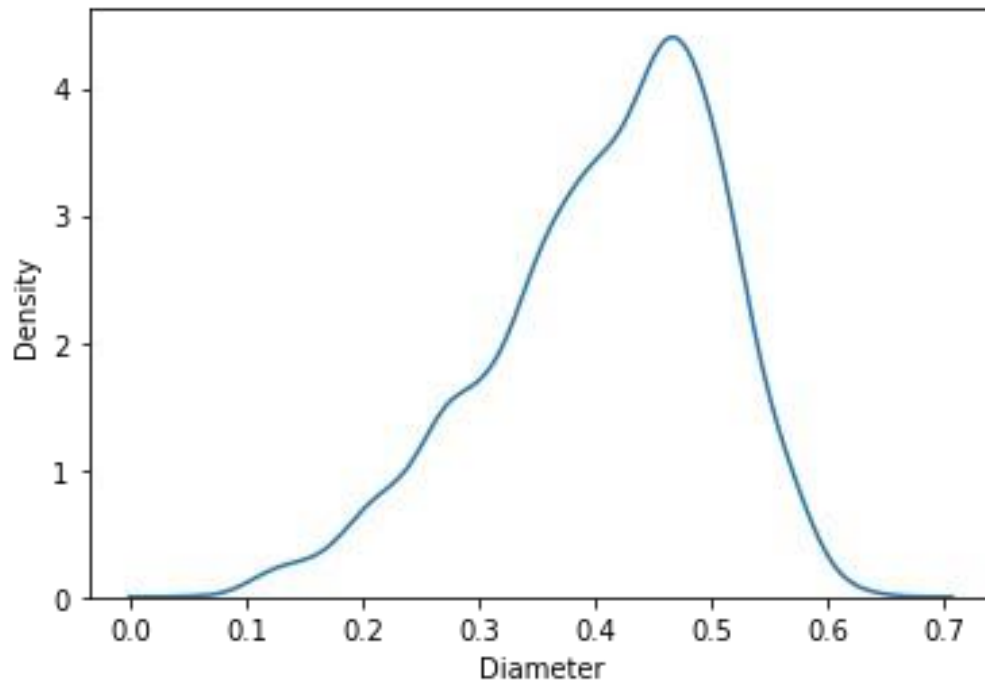
	Shell weight	Rings
0	0.150	15
1	0.070	7
2	0.210	9
3	0.155	10
4	0.055	7

```
a['age']=a['Rings']+1.5
a=a.drop('Rings',axis=1)
```

3.A. univariate Analysis

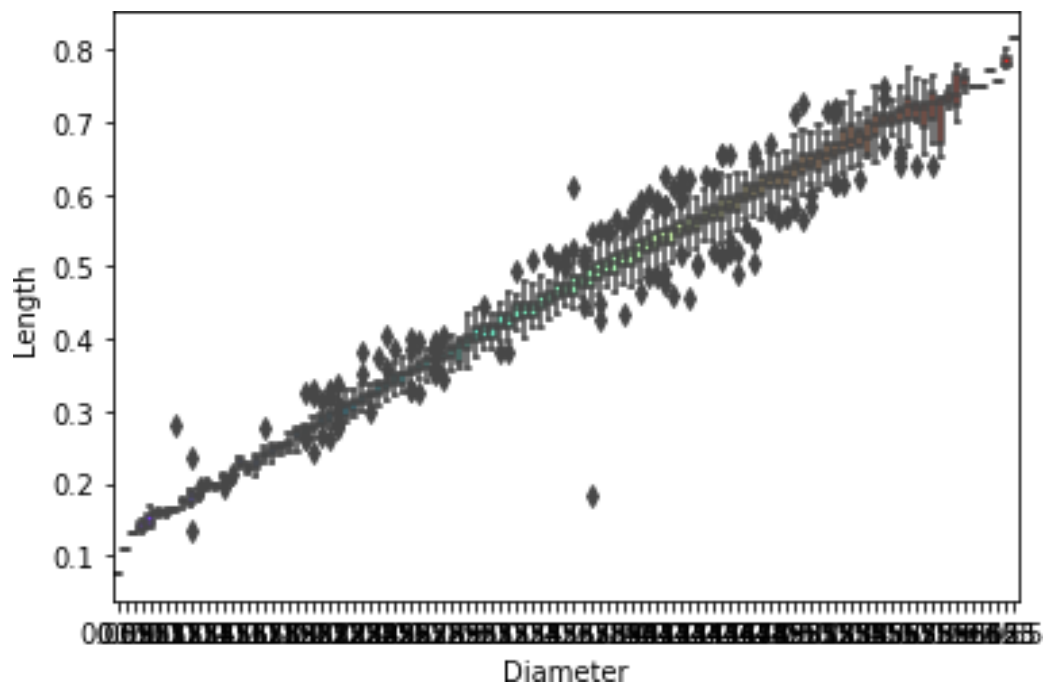
```
ss.kdeplot(a['Diameter'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fbad1961d90>
```



2. Bi-Variate Analysis

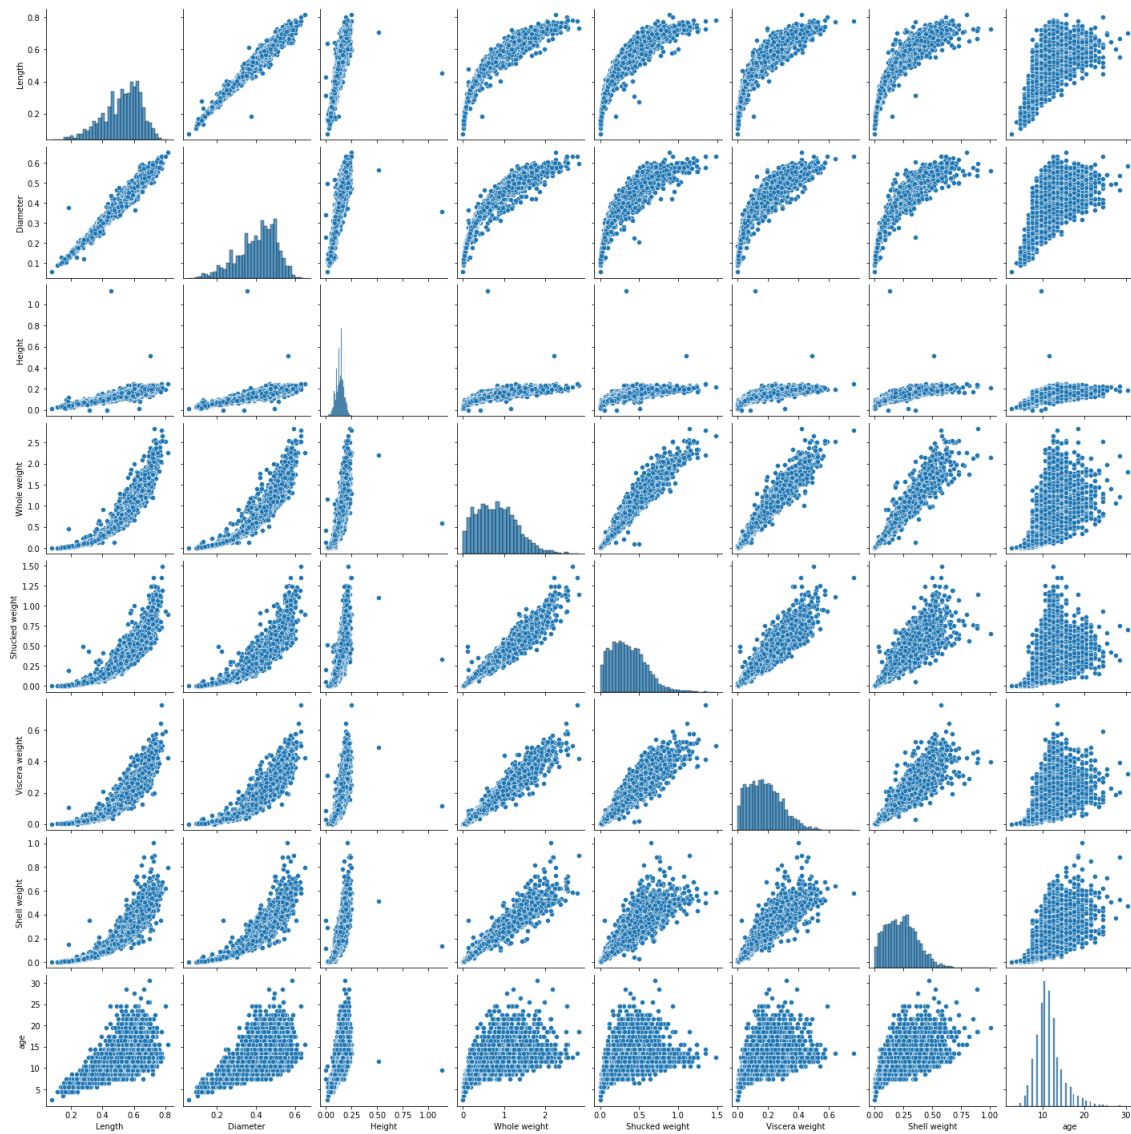
```
ss.boxplot(x=a.Diameter,y = a.Length, palette='rainbow')
<matplotlib.axes._subplots.AxesSubplot at 0x7fbad1844c50>
```



3. Multi-Variate Analysis

```
ss.pairplot(a)
```

```
<seaborn.axisgrid.PairGrid at 0x7fbad096df50>
```



4.Descriptive Statistics

```
a.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 4177 entries, 0 to 4176
```

```
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Sex	4177 non-null	object
1	Length	4177 non-null	float64
2	Diameter	4177 non-null	float64
3	Height	4177 non-null	float64

```
memory usage: 293.8+ KB
```

```
a['Diameter'].describe()
```

```
count      4177.000000
mean        0.407881
std         0.099240
min         0.055000
25%         0.350000
50%         0.425000
75%         0.480000
max         0.650000
Name: Diameter, dtype: float64
```

```
a['Sex'].value_counts
```

```
<bound method IndexOpsMixin.value_counts of 0      M
```

1	M
2	F
3	M
4	I
	• •
4172	F
4173	M
4174	M
4175	F
4176	M

```
Name: Sex, Length: 4177, dtype: object>
```

5. Checking for missing values and dealing with them

```
a.isNull()
```

	Sex	Length	Diameter	Height	Whole weight	Shucked weight \
0	False	False	False	False	False	False
1	False	False	False	False	False	False
2	False	False	False	False	False	False
3	False	False	False	False	False	False
4	False	False	False	False	False	False
...
4172	False	False	False	False	False	False
4173	False	False	False	False	False	False
4174	False	False	False	False	False	False
4175	False	False	False	False	False	False
4176	False	False	False	False	False	False

	Viscera weight	Shell weight	age
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False
...
4172	False	False	False
4173	False	False	False
4174	False	False	False
4175	False	False	False
4176	False	False	False

[4177 rows x 9 columns]

```
a.isnull().sum()
```

```

Length          0
Diameter         0
Height           0
Whole weight     0
Shucked weight   0
Viscera weight   0
Shell weight     0
age              0
Sex_F            0
Sex_I            0
Sex_M            0
dtype: int64

```

6. Find the outliers and replace the outliers


```
a=pd.get_dummies(a)
```

```
dummy_a=a
```

```

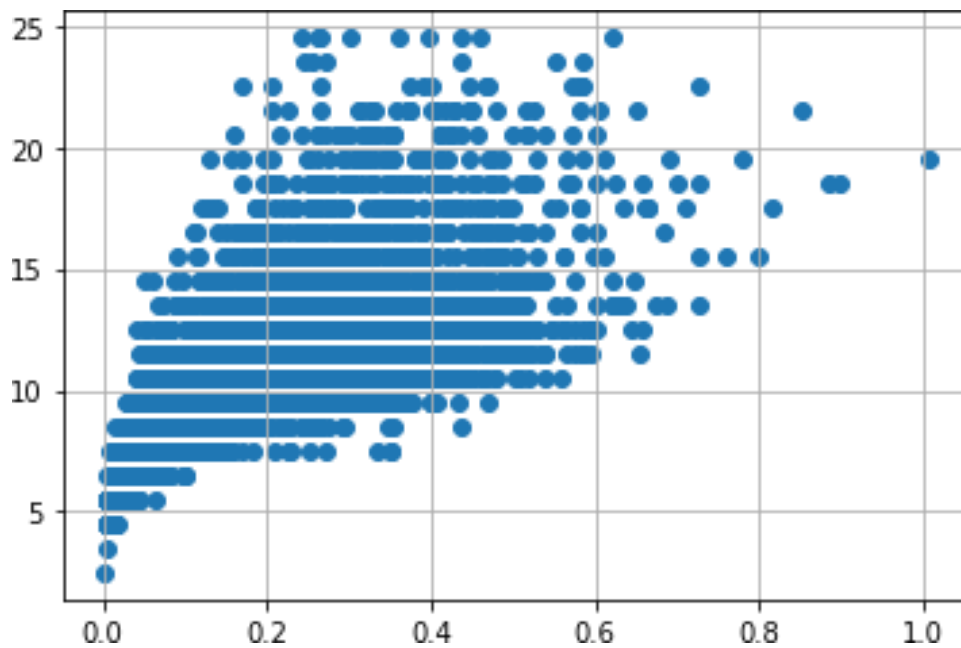
var='Viscera weight'
tlp.scatter(x=a[var],y=a['age'])
tlp.grid(True)

```



```
a.drop(a[(a['Viscera weight']>0.5)&
        (a['age']<20)].index, inplace=True)
a.drop(a[(a['Viscera weight']<0.5)&(a['age']>25)].index, inplace=True)

var='Shell weight'
tlp.scatter(x=a[var],y=a['age'])
tlp.grid(True)
```



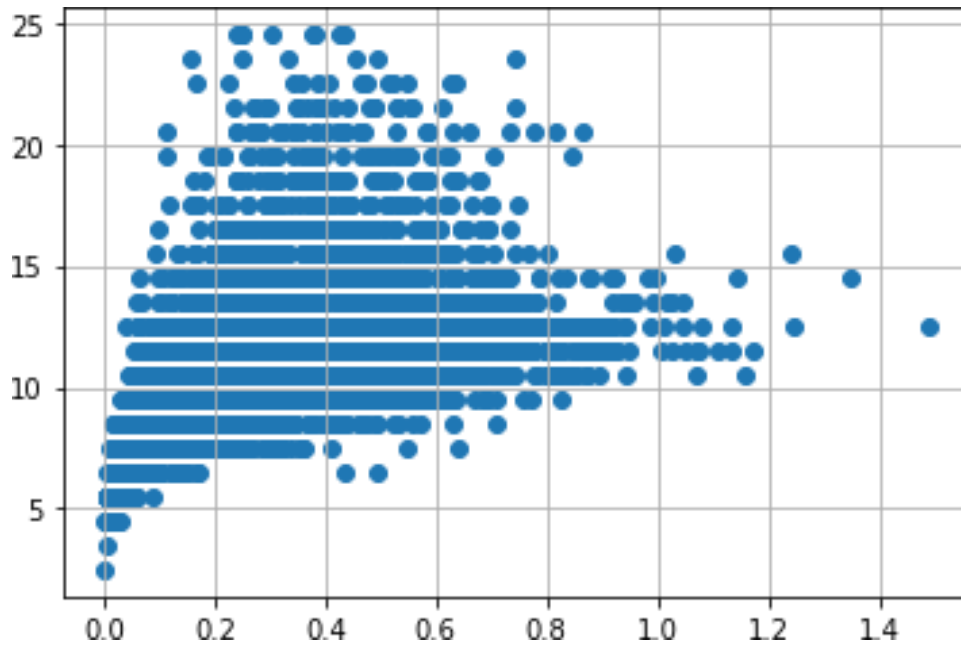
```
a.drop(a[(a['Shell weight'] > 0.6) &
        (a['age'] < 25)].index, inplace = True)
```

```

a.drop(a[(a['Shell weight']<0.8) & (
a['age'] > 25)].index, inplace = True)

var = 'Shucked weight'
tlp.scatter(x = a[var], y =a['age'])
tlp.grid(True)

```

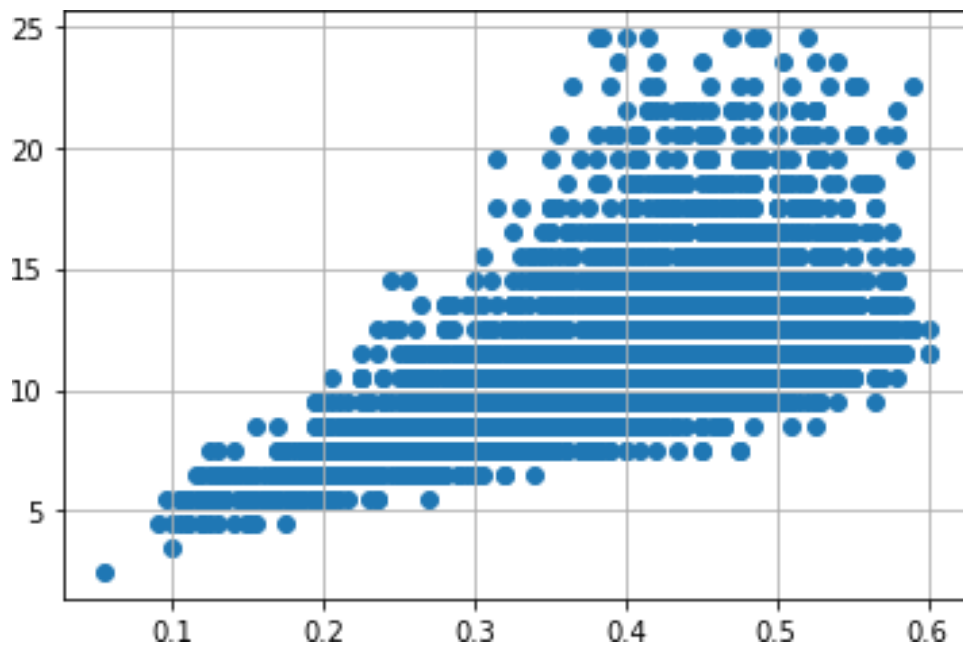


```

a.drop(a[(a['Whole weight'] >= 2.5) &
(a['age'] < 25)].index, inplace = True)
a.drop(a[(a['Whole weight']<2.5) & (
a['age'] > 25)].index, inplace = True)


var = 'Diameter'
tlp.scatter(x = a[var], y = a['age'])
tlp.grid(True)

```



```
a.drop(a[(a['Diameter'] < 0.1) &
          (a['age'] < 5)].index, inplace = True)
a.drop(a[(a['Diameter'] < 0.6) & (
a['age'] > 25)].index, inplace = True)
a.drop(a[(a['Diameter'] >= 0.6) & (
a['age'] < 25)].index, inplace = True)

var = 'Height'
tlp.scatter(x = a[var], y = a['age'])
tlp.grid(True)
```





```

a.drop(a[(a['Height'] > 0.4) &
        (a['age'] < 15)].index, inplace = True)
a.drop(a[(a['Height'] < 0.4) & (
a['age'] > 25)].index, inplace = True)

var = 'Length'
tlp.scatter(x = a[var], y = a['age'])
tlp.grid(True)

```



```

a.drop(a[(a['Length'] < 0.1) &
        (a['age'] < 5)].index, inplace = True)
a.drop(a[(a['Length'] < 0.8) & (
a['age'] > 25)].index, inplace = True)
a.drop(a[(a['Length'] >= 0.8) & (a['age'] < 25)].index, inplace = True)

```

7. Checking for categorical columns

```

numerical_features = a.select_dtypes(include = [np.number]).columns
categorical_features = a.select_dtypes(include = [np.object]).columns

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2:
DeprecationWarning: `np.object` is a deprecated alias for the builtin
`object`. To silence this warning, use `object` by itself. Doing this
will not modify any behavior and is safe.
Deprecated in NumPy 1.20; for more details and guidance:
<https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations>

```
numerical_features
```

```
Index(['Length', 'Diameter', 'Height', 'Whole weight', 'Shucked
weight',
      'Viscera weight', 'Shell weight', 'age', 'Sex_F', 'Sex_I',
      'Sex_M'],
      dtype='object')
```

```
categorical_features
```

```
Index([], dtype='object')
```

Encoding

```
from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
print(a.Length.value_counts())
```

```
0.550    93
0.575    93
0.625    93
0.580    92
0.600    86
..
0.755     2
0.220     2
0.150     1
0.135     1
0.760     1
```

```
Name: Length, Length: 126, dtype: int64
```

```
x=a.iloc[:, :5]
```

```
x
```

	Length	Diameter	Height	Whole weight	Shucked weight
0	0.455	0.365	0.095	0.5140	0.2245
1	0.350	0.265	0.090	0.2255	0.0995
2	0.530	0.420	0.135	0.6770	0.2565
3	0.440	0.365	0.125	0.5160	0.2155
4	0.330	0.255	0.080	0.2050	0.0895
...
4172	0.565	0.450	0.165	0.8870	0.3700
4173	0.590	0.440	0.135	0.9660	0.4390
4174	0.600	0.475	0.205	1.1760	0.5255
4175	0.625	0.485	0.150	1.0945	0.5310
4176	0.710	0.555	0.195	1.9485	0.9455

```
[4096 rows x 5 columns]
```

```
y=a.iloc[:, :5]
```

```
y
```

	Length	Diameter	Height	Whole weight	Shucked weight
0	0.455	0.365	0.095	0.5140	0.2245

1	0.350	0.265	0.090	0.2255	0.0995
2	0.530	0.420	0.135	0.6770	0.2565
3	0.440	0.365	0.125	0.5160	0.2155
4	0.330	0.255	0.080	0.2050	0.0895
...
4172	0.565	0.450	0.165	0.8870	0.3700
4173	0.590	0.440	0.135	0.9660	0.4390
4174	0.600	0.475	0.205	1.1760	0.5255
4175	0.625	0.485	0.150	1.0945	0.5310
4176	0.710	0.555	0.195	1.9485	0.9455

[4096 rows x 5 columns]

9.Spliting the data into training and testing

```
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2)
```

10.Building the model

```
from sklearn.linear_model import LinearRegression
mlr=LinearRegression()
mlr.fit(x_train,y_train)
```

LinearRegression()

11.Training the model

12.Testingthe model

```
x_test[0:5]
```

	Length	Diameter	Height	Whole weight	Shucked weight
2358	0.610	0.485	0.210	1.3445	0.5350
723	0.525	0.410	0.130	0.9900	0.3865
2535	0.640	0.500	0.180	1.4995	0.5930
2717	0.345	0.255	0.095	0.1830	0.0750
29	0.575	0.425	0.140	0.8635	0.3930

```
y_test[0:5]
```

	Length	Diameter	Height	Whole weight	Shucked weight
2358	0.610	0.485	0.210	1.3445	0.5350
723	0.525	0.410	0.130	0.9900	0.3865
2535	0.640	0.500	0.180	1.4995	0.5930
2717	0.345	0.255	0.095	0.1830	0.0750
29	0.575	0.425	0.140	0.8635	0.3930

13.Scaling the independent variables

```

from sklearn.preprocessing import StandardScaler
ss=StandardScaler()
x_train=ss.fit_transform(x_train)

mlrpred=mlr.predict(x_test[0:9])

mlrpred
array([[0.61  , 0.485 , 0.21  , 1.3445, 0.535 ],
       [0.525 , 0.41  , 0.13  , 0.99  , 0.3865],
       [0.64  , 0.5   , 0.18  , 1.4995, 0.593 ],
       [0.345 , 0.255 , 0.095 , 0.183 , 0.075 ],
       [0.575 , 0.425 , 0.14  , 0.8635, 0.393 ],
       [0.57  , 0.48  , 0.18  , 0.9395, 0.399 ],
       [0.61  , 0.485 , 0.165 , 1.087 , 0.4255],
       [0.635 , 0.505 , 0.17  , 1.2635, 0.512 ],
       [0.53  , 0.41  , 0.155 , 0.7155, 0.2805]])

```

14. Measuring the performance using metrics

```

from sklearn.metrics import r2_score
r2_score(mlr.predict(x_test),y_test)

1.0

```